

Introduction

We are required to come up with written-from-scratch algorithms to perform the following computer vision tasks specifically in python:

1. Feature detectors - Identify the interest points
2. Feature descriptors - Extract feature vector descriptor surrounding each interest point
3. Feature matching - Determine correspondence between descriptors in 2 views

Why do we need to learn feature matching and what is the application of it?


To provide you with a clear answer to this question, let's come up with an example. take a look at the following images: This is the Episcopal Palace building in Spain. You can see the same building from different angles and distances.



You as a human can simply identify that there is one unique building in 2 images from 2 different angles and distances. But for a computer, this similarity identification is very difficult and sometimes impossible to find the answer. We are learning Feature Matching to enable a computer to find similar features in different images. Ultimately, the computer will be able to provide the following applications:

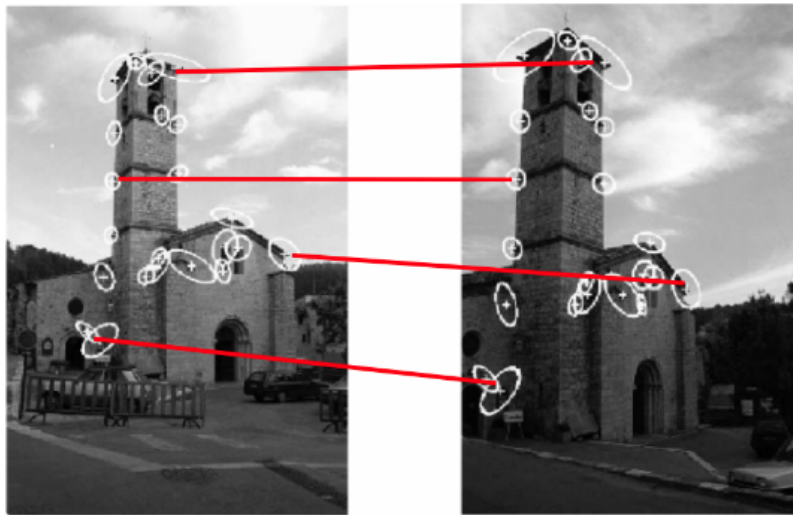
1. Automate object tracking
2. Point matching for computing disparity
3. Stereo calibration
4. Motion-based segmentation
5. Recognition
6. 3D object reconstruction
7. Robot navigation
8. image indexing

In order to write a specific algorithm for feature detection, the followings are required:



Harris Corner Detector

The feature detector or interest point detector algorithm is responsible for finding the most concrete and reliable features in an image that can be easily located in the second image. For instance, the following image pair is illustrated in the following figure:




A good interest point detector should be able to:

1. Detect most of the true interest points
2. deal with noise
3. time-efficient

Mathematical Idea

In this section, we will be using the **Harris Corner detector** algorithm to find interest points in an image. In this algorithm, a small window will be shifted all over the image to find corner points in the image.



Now the question is how to find corner points in an image. To answer this question, we need to dig into the mathematical idea behind Harris Corner Detector:

$$E(u, v) = \sum_{x, y} w(x, y) (I(x+u, y+v) - I(x, y))^2$$

$w(x, y)$: Window Function

$I(x, y)$: Pixel intensity in the image I

$I(x+u, y+v)$: shifted intensity in the image I

We can simply use the Taylor series to find an approximation for $I(x+u, y+v)$:

$$\text{Shifted Intensity} \rightarrow I(x+u, y+v) = I(x, y) + uI_x + vI_y$$

$$E(u, v) = \sum_{x, y} w(x, y) (uI_x + vI_y)^2$$

$$E(u, v) = \sum_{x, y} w(x, y) \left[\begin{bmatrix} u & v \end{bmatrix} \begin{bmatrix} I_x \\ I_y \end{bmatrix} \right]^2$$

$$E(u, v) = \sum_{x, y} w(x, y) \begin{bmatrix} u & v \end{bmatrix} \begin{bmatrix} I_x \\ I_y \end{bmatrix} \begin{bmatrix} I_x & I_y \end{bmatrix} \begin{bmatrix} u \\ v \end{bmatrix}$$

$$E(u, v) = \begin{bmatrix} u & v \end{bmatrix} \left(\sum_{x, y} w(x, y) \begin{bmatrix} I_x \\ I_y \end{bmatrix} \begin{bmatrix} I_x & I_y \end{bmatrix} \right) \begin{bmatrix} u \\ v \end{bmatrix}$$


$$E(u, v) = \begin{bmatrix} u & v \end{bmatrix} M \begin{bmatrix} u \\ v \end{bmatrix}$$

$$M = \sum_{x, y} w(x, y) \begin{bmatrix} I_x I_x & I_x I_y \\ I_x I_y & I_y I_y \end{bmatrix}$$

Now, we need to find the eigenvalues and eigenvectors of the derivative matrix M :

$$\det(M - \lambda I) = 0 \rightarrow \lambda_1 \text{ \& \> } \lambda_2$$

using the eigenvalues for each pixel, we can decide which pixel is the corner point. This concept is illustrated in the following figure:



Having the eigenvalues calculated, we can simply calculate the cornerness of every corner point using the Forstner-Harris equation. with $\alpha = 0.06$. Unlike eigenvalue analysis, this quantity does

not require the use of square roots and yet is still rotationally invariant and also downweights edge-like features.

$$\det(A) - \alpha \times \text{trace}(A) = \lambda_0 \lambda_1 - \alpha (\lambda_0 + \lambda_1)^2$$

Alpha value \rightarrow 0.06.

Extracting feature descriptors

Once the interest points are detected, the next step is finding a correspondence feature descriptor for each interest point. To find feature descriptors, Scale Invariant Feature Transform (**SIFT**) is used.

Now the question is what is SIFT and the idea behind it? Let's answer this question by designing a very straightforward algorithm.

student_sift.py is a python file in which the SIFT algorithm is implemented. get_features(image, x, y, feature_width, scales=None) function receives an input image and its detected interest points.

1. Find the image X(Ix) & Y(Iy) gradient and save them.
2. Use a Gaussian filter to smooth the gradients.
3. Around each interest point(x, y), consider a 16x16 or 32x32 window
4. Calculate the orientation of each pixel. E.g., atan2(Iy/Ix)
5. Extract the histogram of that window. Consider 45 degrees for each histogram bin
6. This Histogram will be your feature descriptor for each interest point
7. Repeat the same process for each interest point and then return all the feature descriptors

Feature Matching

In this section, the K-Nearest Neighbor method is used to find the best match for each feature. Now the question is what is the mathematical idea behind this method. Well, it is very simple to understand and more simple to implement.

Let's say we have 100 features on each image. Here is the algorithm:

1. Pick feature#0 from image#1 and calculate its distance from all features on image#2.
 2. Select the minimum distance and its correspondence indexes on both images. These two indexes match each other by the highest probability.
 3. If that minimum is very near to another one, remove both features from the data.
 4. Repeat the same process for all the features and find the best match for each.
-

Implementation of Harris corner detector

In this section of the report, the Harris corner detector is implemented in Python.

student_harris.py is the file where the algorithm is implemented. `get_interest_points(image, feature_width)` receives an image and a number indicated feature width for finding local maxima.

At the very first step, for the input image, the first gradient in X and Y directions is calculated as follows:

```
1 gaussian_filter_size = 3
2 Gaussian_standard_diviation = 1
3 gaussian_low_pass_filter = cv2.getGaussianKernel(gaussian_filter_size,Gaussian_standard_diviation)
4 img_guassian_filtered = cv2.filter2D(image,-1,gaussian_low_pass_filter)
5 img_grad_y, img_grad_x = np.gradient(img_guassian_filtered)
```

Then the second gradients are calculated as follows:

```
1 img_grad_x_2 = img_grad_x**2
2 img_grad_x_y = img_grad_x*img_grad_y
3 img_grad_y_2 = img_grad_y**2
```

Then the resulted image gradients are convolved with a larger gaussian kernel:



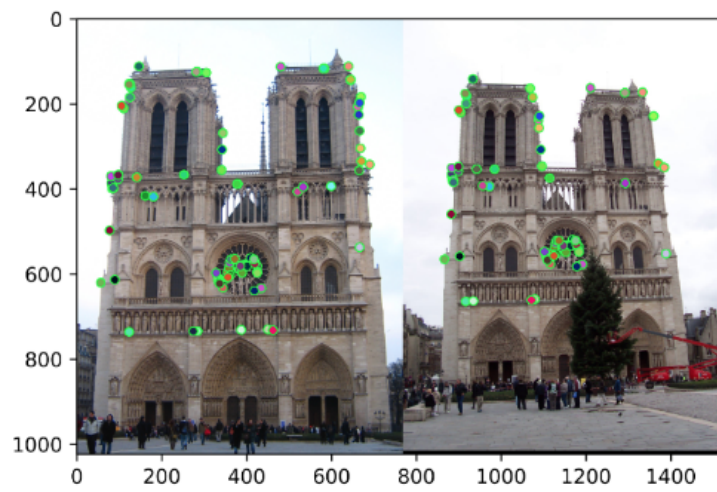
```
1 larger_gaussian_filter_size = 7
2 Gaussian_standard_diviation = 1
3 gaussian_larger_filter = cv2.getGaussianKernel(larger_gaussian_filter_size,Gaussian_standard_diviation)
4 img_grad_x_2_filtered = cv2.filter2D(img_grad_x_2, -1, gaussian_larger_filter)
5 img_grad_x_y_filtered = cv2.filter2D(img_grad_x_y, -1, gaussian_larger_filter)
6 img_grad_y_2_filtered = cv2.filter2D(img_grad_y_2, -1, gaussian_larger_filter)
7
```

Cornerness of the interest points is calculated by Forstner-Harris euation:

```
1 factor_alpha = 0.06
2 det_A = img_grad_x_2_filtered * img_grad_y_2_filtered - img_grad_x_y_filtered **2
3 trace_A = img_grad_x_2_filtered + img_grad_y_2_filtered
4 Forstner_Harris = det_A -factor_alpha * trace_A ** 2
5
6 h, w = Forstner_Harris.shape
7 interest_points = []
8 fh_index = []
9
10 for f1 in range(0, h):
11     fh_index = []
12     for f2 in range(0, w):
13         fh_index = [Forstner_Harris[f1,f2]]
14         fh_index.append(f2)
15         fh_index.append(f1)
16     interest_points.append(fh_index)
```

Given the feature width, the local maxima can be calculated within every feature window. This method is called **Adaptive non-maxima suppression (ANMS)**:

```
1 interest_points_sort_list = sorted(interest_points, reverse = True)
2 number_of_desired_interest_points = 1000
3 interest_points_sort_list = interest_points_sort_list[0:number_of_desired_interest_points]
4 interest_point_neighborhood_list = []
5
6 for i in range(0, number_of_desired_interest_points):
7     local_nibrhd = float('inf')
8     first_location = interest_points_sort_list[i]
9     for j in range(0, i):
10         second_location = interest_points_sort_list[j]
11         euclidean_distance = (first_location[1]-second_location[1])**2 +(first_location[2]-second_location[2])**2
12         local_nibrhd = min(euclidean_distance, local_nibrhd)
13     interest_point_neighborhood_list.append([np.sqrt(local_nibrhd), first_location[0], first_location[1], first_location[2]])
14 x_list = []
15 y_list = []
16 for interest_point_neighborhood in sorted(interest_point_neighborhood_list, reverse = True):
17     x_list.append(interest_point_neighborhood[2])
18     y_list.append(interest_point_neighborhood[3])
19 x = np.array(x_list)
20 y = np.array(y_list)
21
22 return x,y, confidences, scales, orientations
```



Implementation of the feature descriptor

According to the aforementioned algorithm for extracting feature descriptors, the Scale Invariant Feature Transform method is used to find the feature descriptor around each interest point.

At the very first step, the image padding function will be used to pad the input image in order to prepare it for Gaussian filtering:


```

1 def img_padding(img_data, padding_number):
2
3     # if condition: grayscale images
4     if True:
5         i, j= img_data.shape
6         zeros_column = np.zeros((1, i))
7         zeros_row = np.zeros((1, j + 2*padding_number))
8         # for loop: to add zero padding to columns
9         counter = 1
10        for l in range(0, padding_number):
11
12            img_data = np.insert(img_data, 0, zeros_column, axis=1)
13            img_data = np.insert(img_data, j + counter , zeros_column, axis=1)
14            counter = counter + 1
15
16        # for loop: to add zero padding to rows
17        counter = 1
18        for p in range(0, padding_number):
19
20            img_data = np.insert(img_data, 0, zeros_row, axis=0)
21            img_data = np.insert(img_data, i + counter , zeros_row, axis=0)
22            counter = counter + 1
23        return img_data

```

The next step is to write two functions for taking image gradients:

```

1 def get_Ix(inp_img):
2     height, width = inp_img.shape
3     h, w = inp_img.shape
4     Ix = np.zeros((h-1, w-1))
5     for row in range(0, h-1):
6         for col in range(0, w-1):
7             Ix[row, col] = inp_img[row, col + 1] - inp_img[row, col]
8     return Ix
9
10 def get_Iy(inp_img):
11     h, w = inp_img.shape
12     Iy = np.zeros((h-1, w-1))
13     for col in range(0, w-1):
14         for row in range(0, h-1):
15             Iy[row, col] = inp_img[row + 1, col] - inp_img[row, col]
16     return Iy

```

Now, we are ready to dig into the main function of the feature descriptor:

Don't Panic! 😊 I wrote this code very simple so that everyone can make sense of it. After taking the X and Y derivatives of the image, we will define the bins of the histogram. Then we will start looping over the image, stopping at each interest point, and then extract the feature descriptor, histogram of orientation and magnitude. This way the descriptor vector for each interest point is extracted and we return them in normalized shapes.


```

1 def get_features(image, x, y, feature_width, scales=None):
2     padding_number = int(abs(feature_width/2))
3     image = img_padding(image, padding_number)
4     dx = get_Ix(image)
5     dy = get_Iy(image)
6     features_descriptors=[]
7     number_of_interest_points = x.size
8     bin_val_1 = -1
9     bin_val_2 = 1
10    bin_val_3 = 2
11    for interest_point in range(0, number_of_interest_points):
12        o = int(abs(feature_width/4))
13        pixel_magintude_all = np.zeros((o,o,8))
14        for j in range(0, feature_width):
15            for i in range(0, feature_width):
16                ix_value = dx[(int)(y[interest_point]) + j, (int)(x[interest_point]) +i]
17                iy_value = dy[(int)(y[interest_point]) +j, (int)(x[interest_point]) +i]
18                pixel_angle = np.arctan2(iy_value, ix_value)
19                if pixel_angle > bin_val_2:
20                    pixel_angle = bin_val_3
21                if pixel_angle < bin_val_1:
22                    pixel_angle = bin_val_1
23                pixel_magintude = np.sqrt(ix_value**2 + iy_value**2)
24                if ix_value >0:
25                    pixel_magintude_all[(int)(j/4), (int)(i/4), math.ceil(pixel_angle+1)] += pixel_magi
26                else:
27                    pixel_magintude_all[(int)(j/4), (int)(i/4), math.ceil(pixel_angle+5)] += pixel_magi
28            single_feat = np.reshape(pixel_magintude_all,(1,o*o*8))
29            single_feat = single_feat/(single_feat.sum())
30            features_descriptors.append(single_feat)
31    fv = np.array(features_descriptors)
32    return fv

```

Implementation of feature matching

The feature matching is our last algorithm to implement in this assignment. According to the aforementioned algorithm for feature matching:

1. start looping over all the returned feature descriptors.
2. Pick the feature descriptor from feature1 array and compare them with all feature descriptor from feature2.
3. Detect the first and second minima. If they are very close to each other, ignore this interest points. Otherwise, save the indexes.
4. calculate the confidence by dividing the first and second minima in each feature descriptor.
5. Repeat the same process for all feature descriptors.
6. return matches and confidence.

```

1 def match_features(features1, features2, x1, y1, x2, y2):
2     d_all = []
3     feat_1_match_index = []
4     feat_2_match_index = []
5     feat_match = []
6     mathing_confidence = []
7
8     for i in range(features1.shape[0]):
9         for j in range(features2.shape[0]):
10             local_distance = np.linalg.norm(features1[i]-features2[j])
11             d_all.append(local_distance)
12             first_min = min(d_all)
13             min_1_index = d_all.index(first_min)
14             del d_all[min_1_index]
15             second_min = min(d_all)
16             min_2_index = d_all.index(second_min)
17             nearest_neighbor_distance_ratio = first_min /second_min
18             match_index = min_1_index
19             feature_index_1 = i
20             feat_1_match_index.append(i)
21             feat_2_match_index.append(match_index)
22             mathing_confidence.append(nearest_neighbor_distance_ratio)
23             feat_match.append([nearest_neighbor_distance_ratio, feature_index_1, match_index])
24             d_all = []
25
26     match_list = []
27     confidence_list = []
28     for mt in sorted(feat_match)[:100]:
29         match_list.append([mt[1], mt[2]])
30         confidence_list.append(mt[0])
31
32     matches = np.array(match_list)
33     confidences = np.array(confidence_list)
34
35     matches = matches.astype(int)
36     return matches, confidences

```

Harris corner detection - Results

The output of the aforementioned algorithm is the XY location of all interest points within the processed image. The following image pairs are processed:

Notre Dame Building: accuracy

Settings to get the best results:

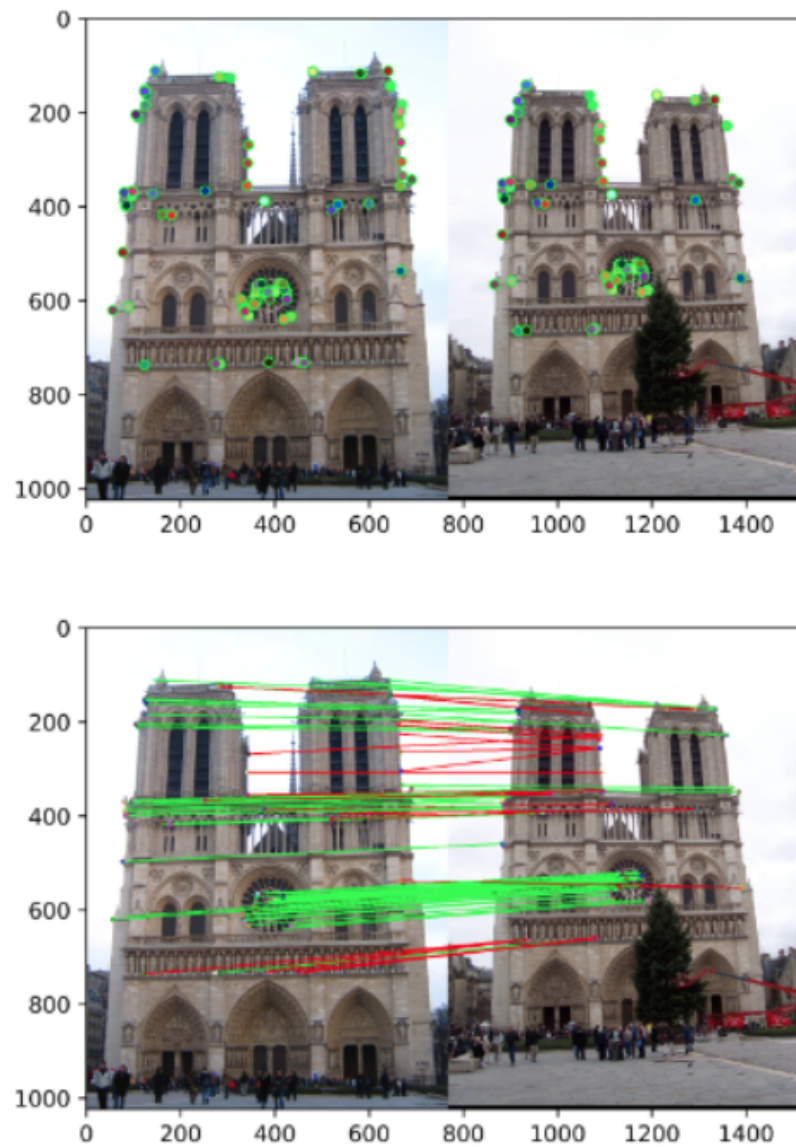
Feature Width: 32

Alpha = 0.06

Gaussian Distribution 3x3 → Standard Deviation: 1.0 → applied on the input image

Gaussian Distribution 7*7 → Standard Deviation: 7.0 → applied on image gradients

Accuracy: 73%



Mount Rushmore:

Settings to get the best results:

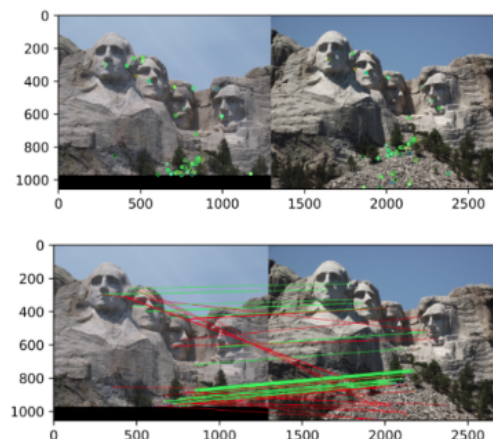
Feature Width: 32

Alpha = 0.005

Gaussian Distribution 3x3 → Standard Deviation: 1.0 → applied on the input image

Gaussian Distribution 7*7 → Standard Deviation: 10 → applied on image gradients

Accuracy: 61%



Episcopal Gaudi:

Settings to get the best results:

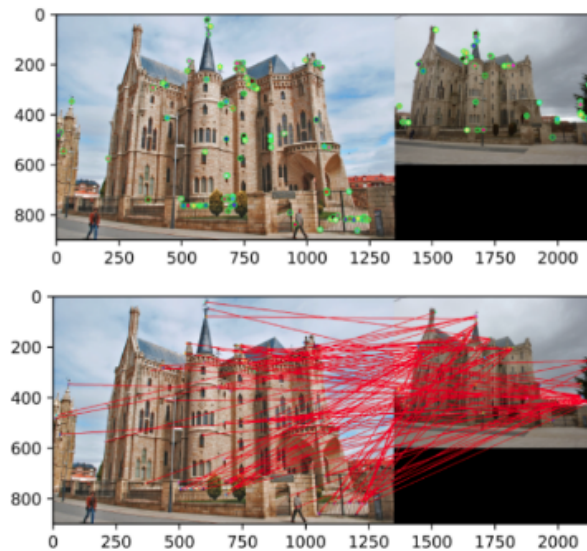
Feature Width: 32

Alpha = 0.005

Gaussian Distribution 3x3 → Standard Deviation: 1.0 → applied on the input image

Gaussian Distribution 7*7 → Standard Deviation: 10 → applied on image gradients

Accuracy: 0.0 %



References

[1]: Richard Szeliski. 2010. Computer Vision: Algorithms and Applications (1st. ed.). Springer-Verlag, Berlin, Heidelberg.

[2]. Abdel-Hakim, A. E. and Farag, A. A. (2006). CSIFT: A SIFT descriptor with color invariant characteristics. In IEEE Computer Society Conference on Computer Vision and Pattern Recognition (CVPR'2006), pp. 1978– 1983, New York City, NY.

[3]. Adelson, E. H., Simoncelli, E., and Hingorani, R. (1987). Orthogonal pyramid transforms for image coding. In SPIE Vol. 845, Visual Communications and Image Processing II, pp. 50–58, Cambridge, Massachusetts.